

# Reprovide Sweep

*Opening the DHT to large content providers*



Gui Michel  
*@guissou*

Probelab,  
Protocol Labs

IPFS Thing  
16th April 2023



Protocol  
Labs

# DHT Provide Process

When a node wants to advertise to the DHT that it provides CID

- 🟡 DHT lookup to find the `rep1` (20) closest peers

  - 🟢 Number of hops:  $\sim 3$

  - 🟢 Number of inflight messages:  $\sim 10$

  - 🟢 Total number of sent messages:  $\sim 30$

tes

- 🟡 Allocate the Provider Record to these `rep1` closest peers

# DHT Provide Process

- Number of connections opened per provide:  $\sim 35$ 
  - Assume that we already have an open connection for peers in the first hop
  - Intermediary lookup peers:  $\sim 15$
  - rep1 closest peers: 20
- Number of messages sent per provide:  $\sim 70$ 
  - DHT Lookup:  $\sim 30 + 20$
  - Provide messages: 20

# Large Content Providers

Large content provider reproviding 1 Billion CIDs every `ReprovideInterval` (22h)

- Number of opened connections:  $\sim 35B \approx 450'000/s$
- Number of messages sent:  $\sim 70B \approx 900'000/s$

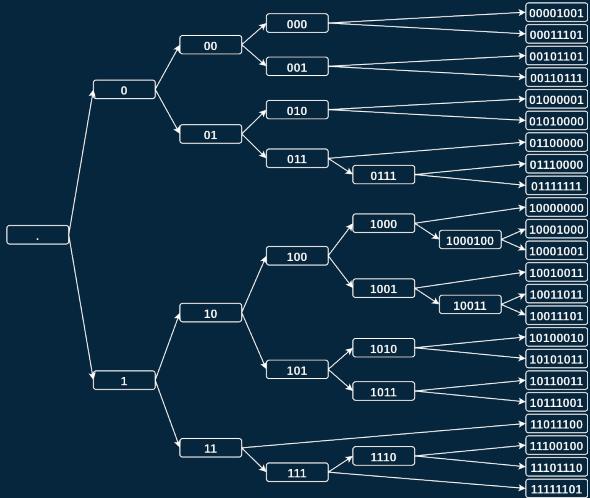
# Optimizing the Reprovide operation

- More CIDs to reprovide than DHT Servers
  - $\Rightarrow$  multiple CIDs are allocated on the same DHT Server
  - Group CIDs allocated on the same DHT Server and reprovide them sequentially
  - Periodically sweep the keyspace and reprovide CIDs
- 
- Minimizing the number of connections to open and messages to send
  - Number of connections to open  $\approx$  number of DHT server peers



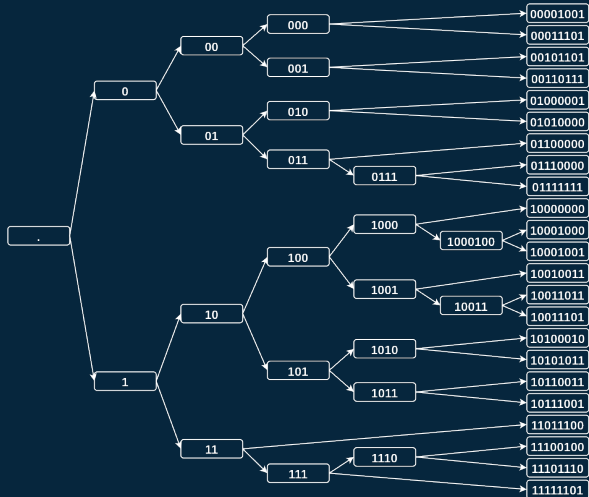
# Binary Trie

- Binary Prefix Tree
- Data structure optimized for working with XOR distance
- Helps visualize locality in the Kademlia keyspace
- We will use them to group *close* CIDs



# Keyspace Regions

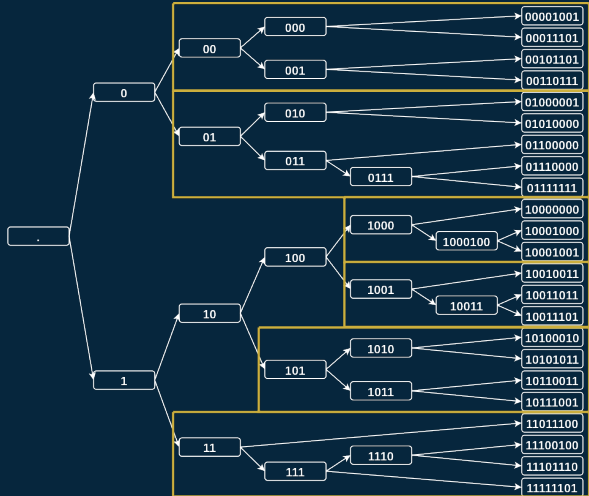
- ⬢ A Region is defined as a prefix of the **peers** binary trie that has at least `rep1` leaves.
- ⬢ We are interested in the smallest possible regions fully covering the keyspace



# Keyspace Regions: Example

Example:  $\text{rep1} = 3$

- Keys represent peers
- A CID is only stored in the region matching its prefix.





# Region **exploration**

- 🔹 Lookup a random key within the target Region
- 🔹 If some returned peers don't match the Region's prefix, Region is fully explored
- 🔹 Else, take the largest fully explored subregion, and explore its neighbor subregion
- 🔹 Repeat until the Region is fully explored

# Sweep

- ✧ Explore Kademlia keyspace `Regions` to sweep the keyspace *from left to right*
- ✧ All `Provider Records` belonging to the looked up `Region` are reprovided
- ✧ The keyspace sweep takes `ReprovideInterval` to complete
- ✧ Peers are stored in a binary trie, in order to define the `Regions`
- ✧ CIDs are stored in a distinct binary trie, for fast sweep, insert and delete

# Reproviding for a Region

- Define a temp key-value store  $\text{PeerID} \rightarrow [\text{Keys}]$  for PeerIDs within the Region
- For all keys  $k$  belonging to a Region, add  $k$  to its `rep1` closest peers
- Iterate over all PeerIDs within the Region and reprovide all associated keys
- The number of workers can easily be limited

# Region **Reprovide** Scheduling

- Each Region should be republished within `ReprovideInterval`  $\pm$  small delay
- Once a sweep cycle, the delays at which the Regions are reprovided are adjusted
- The Scheduler keeps track of the time the last reprovide happened for all Regions
- Enable resuming reprovide quickly if the node goes offline for some time

# First Provide

- ⬢ First Provide is timely, provide immediately
- ⬢ Reprovide isn't timely
- ⬢ The first Reprovide of a CID is likely to happens less than `ReprovideInterval` after its first provide

# Region **shrinking**

- ⬢ A Region can shrink in the number of peers from one exploration to the next one
- ⬢ A Region containing less than `rep1` peers must be merged with its closest neighbor
- ⬢ Providers Records are republished at the planned time, but the scheduler reschedule the reprovide time for the next round

# Region **expansion**

- ⬠ A Region can grow in the number of peers from one exploration to the next one
- ⬠ A Region can be split in two Regions if both of its branches have at least `repl` peers
- ⬠ Regions must always be splitted when possible
- ⬠ When a Region is split in two (or more), the Provider Records associated with the new Regions are reprovided concurrently
- ⬠ The scheduler is responsible to space the reprovide time for the next round

# Intuition for proof of correctness

- ✦ First provides in appropriate Region, Region added to scheduler
- ✦ When a new CID must be provided, it is added to the appropriate Region
- ✦ Regions can be split or combined during peer churn within a Region



# Performance Evaluation

Reproviding 1B CIDs to a 20'000 peers network using  $repl=20$ .

- Number of Regions:  $\sim \frac{\#peers}{2 \times repl} \approx 500$
- Number of connections opened to explore a Region:  $\sim 55$
- Number of messages sent to explore a Region:  $\sim 70$
- Total number of connections opened:  $\sim 500 \times 55 \approx 28'000$ 
  - Vanilla Provide:  $\sim 35B$ , improvement:  $\sim 1M\times$
- Total number of messages sent:  $\sim 500 \times 70 + 20B \approx 20B$ 
  - Vanilla Provide:  $\sim 70B$ , improvement:  $3.5\times$

# Comparison with ProvideMany from the Accelerated DHT Client

Accelerated DHT Client ProvideMany:

- + Group CIDs by XOR distance before reproviding
- All CIDs are reprovided at the same time → rush hour
- No DHT lookup before providing → routing table stale entries
- Running a crawler to refresh the routing table → expensive
- Constant CIDs groups size → additional messages and connections

# Migrating the Reprovide responsibility to the Content Routers

- ✦ The current DHT implementations only expose a `Provide(CID)` interface
- ✦ Currently IPFS implementations must handle the reprovide
- ✦ Different Content Routers have different reprovide mechanisms
- ✦ Each Content Router should be responsible to reprovide content
- ✦ Interface should be: `ProvideContent(CID)`, `UnprovideContent(CID)`, `GetProvidedContent()`

# Conclusion

- ⬢ Minimize Reprovide cost for everyone!
- ⬢ Enable large Content Providers to use the DHT
- ⬢ Once everyone uses the DHT, we can start moving away from the Bitswap broadcast
- ⬢ To be shipped with the Double Hash DHT later this year



Reprovide Sweep Notion Page